

PPC and Nurture: Executive Summary

Version 9 | 2 June 2026

The one-page version

The problem. We inherited the paid accounts from three freelancers, across Google, LinkedIn, Microsoft and Meta, in the UK, Iberia and France. They are a sprawl: broken tracking, money going to the wrong products in the wrong markets, narrow audiences that never convert, and bid strategies left on autopilot. Nobody could run them well, and the reporting undercounts what paid actually brings in.

The plan, in a line. Tear it down and rebuild lean. Roughly 33 focused campaigns instead of the sprawl, one flagship product per market, each funded properly, named so it reports cleanly, and pointed at a real conversion.

Paid is only half of it. The other half is nurture, what happens to a lead after it comes in. Today most leads dead-end. We fix that, so every lead gets worked and the demand we pay for turns into pipeline. For a slow, considered B2B purchase like ours, that second half is where the deal is won.

What it needs from the wider team. Most of the remediation is ours to execute, the four ad platforms and the tag container, with the website-migration lead, because the conversion outage is a site dataLayer break rather than an attribution one. A short set genuinely sits with RevOps (the cross-system attribution into Dynamics and the offline-conversion connection), and the lead-routing engine RevOps built is already live. Until the conversion signal is firing again, campaigns still launch and still report through naming and UTMs; they just can't optimise toward conversions yet. No new budget: we fund the rebuild by redeploying the waste we found. The committed recovery is €35-55K a year, planned as a conservative and an optimistic case (the identified waste annualises to roughly €105K at the current run-rate; €35-55K is the realistic net we hold ourselves to).

Paid captures the lead > Score + nurture > MQL > Sales accepts (SAL) > SQL and pipeline

Budget (no new spend)

Funded by redeploying identified waste. Planning ranges, to firm up once we have a reliable cost per lead by channel.

PLATFORM	UK	IBERIA	FRANCE	LATAM	MONTHLY
Google	€4-6K	€4-6K	€4-6K	€0.6-1K	~€14-19K
LinkedIn (70/30)	€3-4K	€3-4K	€1.5-2K	-	~€7.5-10K
Microsoft	€1-2K	Q3 build	€0.8-1.6K	incl.	~€2-3.5K
Meta (75/25)	€0.5-1K	€0.5-1K	€1-1.5K	-	~€2-3.5K

~33 campaigns total (Google ~13, LinkedIn ~9, Microsoft ~4, Meta ~7). One person can run weekly hygiene across 33 focused campaigns; the current sprawl can't be run well by anyone.

How we measure success

Judged on pipeline, not platform-reported clicks. HubSpot pipeline, closed-loop through the offline conversion upload, is the scoreboard.

METRIC	ROLE
Cost per lead, by channel and market	Launch-phase efficiency gate (native social forms ran €26-55; ceilings €60-90 social, €120-150 search)
MQL-to-SQL rate	Lead quality, matters more than volume
Cost per SQL	True acquisition efficiency, replaces cost-per-click as the headline
Pipeline and win rate	Revenue contribution (paid's trailing-12-month win rate is 15.8%, the baseline to hold or beat)
SAL-to-booking (nurture)	From sub-15% toward 20% by end of 2026

Value over volume. Optimise for lead value, not count, by pushing values back from HubSpot (MQL €500, SQL €2,000, Opportunity €10,000, Closed Won €25,000).

Measure lift, not last-click. A brand-defence geo holdout at launch (hold brand out in matched regions over a fixed window, then compare), because cheap last-click spend can hide demand that would convert anyway.

Reallocate, don't use-it-or-lose-it. Budget moves monthly toward the best cost per SQL; read year-on-year because hospitality demand is seasonal; plan a conservative and an optimistic case, not one forecast.

Full build detail, per-platform specs and the tracking and routing remediation are in the plan that follows.

PPC and Nurture: The Plan

Version 9 | 2 June 2026 | Companion: the per-platform execution document

WHAT CHANGED IN V9

Re-cut the tracking work by real ownership and folded in the live routing engine. The conversion outage is now framed as a site dataLayer break owned by us and the website-migration lead, not a RevOps fix; section 14 is reorganised around who owns each item, with a short genuine RevOps list (the cross-system attribution into Dynamics and the offline-conversion connection). Non-demo lead routing is live (RevOps), so the routing gap is reframed to the residual, a working signal on non-demo leads. Automatic lead scoring is RevOps Phase 2. Paid suppression and the fit gate now use the under-20-employee company-size floor, matching the inbound qualification gate.

WHAT CHANGED IN V8

Gave the expansion data path an owner and a date: product-usage data in the CRM (RevOps) and the customer-health score exposed to marketing (Customer Success) are explicit owned items in the section 14 list, both targeted Q3, gating cross-sell only, not the acquisition rebuild.

The one-page version

The problem. We inherited the paid accounts from three freelancers, across Google, LinkedIn, Microsoft and Meta, in the UK, Iberia and France. They are a sprawl. Broken tracking, money going to the wrong products in the wrong markets, narrow audiences that never convert, and bid strategies left on autopilot. Nobody could run them well, and the reporting undercounts what paid actually brings in.

The plan, in a line. Tear it down and rebuild lean. Roughly 33 focused campaigns instead of the sprawl, one flagship product per market, each one funded properly, named so it reports cleanly, and pointed at a real conversion. We rebuild on a clean account rather than tidying the mess.

Paid is only half of it. The other half is nurture, what happens to a lead after it comes in. Today most leads dead-end. We fix that, so every lead gets worked and the demand we pay for actually turns into pipeline. For a considered, slow B2B purchase like ours, that second half is where the deal is won.

What it needs from the wider team. Most of the remediation is ours to execute, the four ad platforms and the tag container, with the website-migration lead, because the conversion outage is a website dataLayer break rather than an attribution one (the full ownership map is in section 14). A short set genuinely needs RevOps, the cross-system attribution into Dynamics and the offline-conversion connection, and the lead-routing engine RevOps built is already live. Until the conversion signal is firing again, campaigns still launch and still report through naming and UTMs. They just can't optimise toward conversions yet. There is no new budget: we fund the rebuild by redeploying the waste we found. The committed recovery is €35-55K a year, planned as a conservative and an optimistic case (the identified waste annualises to roughly €105K at the current run-rate, €67K is the post-haircut realistic recovery, and €35-55K is the net we commit to; the report carries the full methodology).

How to read this. The summary above is the whole story. Sections 1 to 8 are the paid build. Section 9 is the nurture engine. Section 10 is the expansion engine, cross-sell and up-sell into the customer base. Sections 11 and 12 cover creative and how we measure success. Sections 13 to 15 are budget, the tracking and routing remediation (who owns what), and the launch order. Skip to whatever is yours.

Paid captures the lead › Score + nurture › MQL › Sales accepts (SAL) › SQL and pipeline

The whole engine in one line. Paid remarketing runs alongside it, keeping the brand present between nurture emails. Today most non-demo leads die at the "score + nurture" step; fixing that is the heart of the plan.

Contents

1. How to use this with the execution doc
2. The rules every campaign follows
3. The product-market map (the filter on everything)
4. The teardown: what we pause and archive
5. Google Ads, by market
6. LinkedIn Ads, by market
7. Microsoft Ads, by market
8. Meta Ads, by market
9. The nurture engine paid feeds into
10. The expansion engine: cross-sell and up-sell
11. Creative and testing
12. How we measure success
13. Budget and campaign count
14. Tracking and routing remediation (who owns each)
15. The launch order

Appendix A. ABM build detail (folded in from the former standalone spec).

1. How to use this with the execution doc

Short version: read this for the why, build from the execution doc.

This document is the thinking: what we run, why it exists, who it targets, and the settings behind it. The execution doc is the doing: step-by-step blocks per platform and market. Build from that, adapting each block to what you see on screen rather than running it all blind. If the two ever disagree, this document wins.

One thing worth saying plainly about the method. For every account the order is the same: export the current state, pause everything, archive it under a dated label, then build fresh. We don't edit old campaigns into new ones; section 4 explains why.

2. The rules every campaign follows

Short version: a few rules every campaign obeys, so nothing launches half-built and everything reports cleanly.

These apply to all four platforms. Miss one and the campaign doesn't go live.

Naming. The platform campaign name and the utm_campaign value are the same string, so the data lines up in both Dynamics and HubSpot. The prefix matches the Dynamics campaign theme; the suffix adds the PPC detail.

LEVEL	PATTERN
Campaign	{REGION}-{ICP}-{PRODUCT}-{OBJECTIVE}-{THEME}-{QYY}
Ad set / ad group	{Audience}-{Geo}-{Lang}-{FunnelStage}
Ad / creative	{Format}-{Angle}-V##

SLOT	ALLOWED VALUES
REGION	UK, ES, FR, LATAM, GLOBAL
ICP	QSR, FSR, QSRFSR, HOTEL, CC (contract catering), HC (healthcare), KEYACC, ALL
PRODUCT	Flow, Workforce, Easilys, Compliance, Engage, MapalOS, Brand
OBJECTIVE	Brand, Search, LeadGen, Demand, Retarget, ABM
THEME	kebab-case, max 4 words, matches the Dynamics theme
QYY	Q226 (quarter and year)

Hyphens only, no spaces, no underscores inside a slot, 80 characters max. Nobody types these by hand; they come from the builder sheet, so the name and the UTM always match. Example: UK-HOTEL-Flow-Search-Always-On-Q226.

UTMs, and where they land in HubSpot.

PARAMETER	VALUE	LANDS IN HUBSPOT AS
utm_source	google / linkedin / bing / meta	Latest traffic source
utm_medium	cpc / paid-social / display / video	Latest traffic source
utm_campaign	full campaign name (same as platform)	Drill-down + last-touch campaign
utm_content	adset-name _ad-name	Drill-down 2
utm_term	keyword (search only)	Drill-down detail

Why naming is sacred

LinkedIn Lead Gen Forms and Meta Instant Forms don't pass UTMs. On those campaigns the campaign name is the only thing HubSpot receives, so a wrong name means an unattributable lead. That's why the builder sheet governs every name.

Conversions. Every campaign points at the right conversion from day one, even where that conversion isn't firing cleanly yet. So the moment the conversion signal is firing cleanly again, the campaign already has the right goal and needs no rework. Three conversions, in funnel order: Lead MQL (the one we optimise to), Lead SQL, and Demo Form Submission. The HubSpot attribution setup itself is sound; the conversion break is upstream of it, in the site dataLayer and the tag container.

Tracking prerequisite (ours, with the website lead). The conversion signal is down because the website migration changed the dataLayer contract these conversions depend on: form conversions were moved onto a HubSpot form event and the form-id keys were then renamed, so the triggers stopped matching. The Google Ads conversions themselves are live and correctly configured (the FR Lead MQL conversion and the per-account demo conversions); they are simply not being triggered, and ES and UK Lead MQL conversions don't exist yet and are part of the build. Realigning that contract sits with us (the tag container) and the website-migration lead, not RevOps. A main Meta pixel is also dead, ours to restore. Until the signal fires, campaigns launch and report via naming and UTMs but can't optimise toward conversions. Full ownership map in section 14.

Bidding: start honest, then switch. Because the conversion signal isn't reliable yet, every campaign launches on an interim bid strategy that's honest about the data we have, then switches to conversion-based bidding about seven days after the conversion signal is confirmed firing.

CAMPAIGN TYPE	AT LAUNCH	AFTER THE TAG FIX
Brand defence	Manual CPC or impression share	Stays manual. Never Maximize Clicks.
Flagship search	Maximize Clicks with a firm CPC cap	Maximize Conversions, then Target CPA at 30+ conversions/month
Lead form (social)	Maximize leads / lowest cost	Cost cap once we have a cost-per-lead baseline
Remarketing	Manual or cost cap, small budget	Conversion-based once audiences are populated

The interim phase is not set-and-forget. It's a data-gathering phase with a defined exit. We don't leave Maximize Clicks running on a flagship forever, which is exactly the failure the audit found across ES, UK and FR.

Interim guardrails (set at launch, not later). While bidding runs blind, it gets hard limits from day one. Starting values, to be tuned with live cost-per-lead data, but never left unset.

GUARDRAIL	VALUE	IF IT BREACHES
Max CPC, brand	UK £2-3 / ES, FR €2-3	Brand should be cheap. If CPC climbs, check for competitor encroachment, don't just raise the cap.
Max CPC, flagship search	€6-9 (by market)	Pause the ad group, check search terms and Quality Score before lifting.
Cost-per-lead ceiling, search	€120-150	Pause and diagnose (audience, landing page, offer).
Cost-per-lead ceiling, social	€60-90 (native forms ran €26-55 in the audit)	Pause, check audience size and form friction.
Switch-to-conversion threshold	30+ Lead MQLs per campaign / 30 days	Below this, stay interim or use value-based bidding (below).

Value-based bidding (the quality lever). Ours is a considered, low-volume purchase. Optimising for lead volume just rewards cheap, low-intent leads. So instead we optimise for lead value: we push a value back to the platform as each lead moves through the funnel, and the platform learns to find the kind of person who becomes a real deal. It works even when a campaign can't hit 30 conversions a month.

STAGE	VALUE PUSHED TO THE PLATFORM
MQL	€500
SQL	€2,000
Opportunity	€10,000
Closed Won	€25,000

This runs on Google and LinkedIn once the prerequisites are met, and it's the default for the flagship and lead-gen campaigns, ahead of plain volume bidding. The prerequisites: offline conversion upload switched on, MQL set automatically (not by a rep flipping a stage by hand, which misses the match window), and the LinkedIn click ID plus UTM's captured on every form. All three are in section 14.

Attribution windows. Senior hospitality buyers don't convert in one session, so we set a 30-day click-through window across platforms, with view-through handled conservatively. Windows that are too short undercount paid and feed the eventual smart-bidding switch bad data, so we get this right before trusting any conversion-based bidding.

Audiences: broad but correct, and who we exclude. No always-on social audience runs below 50,000, the platform's delivery threshold. Narrow hand-built lists starve the algorithm, which the audit proved on LinkedIn (a 3,400 list produced zero leads against a 22,000 audience that worked). Narrow, named-account targeting moves to sales sequencing, not paid.

Suppression, on every campaign. Quality is protected by who we leave out as much as who we target. From day one, every campaign excludes:

- Existing MQLs and above (synced from HubSpot, refreshed weekly), so we don't pay to re-touch people already in the funnel.
- Job titles Student, Intern, Assistant (the standalone word), which carry no buying authority.
- Companies under 20 employees, matching the inbound qualification gate, which catches list mismatches and sub-ICP traffic.
- Current customers and open opportunities (they belong in cross-sell and sales motions, not acquisition).
- Clearly wrong verticals, per the product-market map.

The audience layers beyond cold targeting. Once HubSpot holds clean segments, paid becomes a multiplier on the database we already own, not just a way to find strangers. Three layers sit on top of cold targeting. We specify the data pipeline now and switch them on behind the tracking fix.

LAYER	WHAT IT IS
Database activation	Ads to HubSpot segments via LinkedIn Matched Audiences and Google Customer Match, with different messaging by funnel depth (cold sees the problem, warm sees product and proof, MQL sees a demo CTA).
Lookalikes from closed-won	LinkedIn Lookalikes and Google Similar Audiences seeded from closed-won by product and ICP. Beats job-title targeting once the seed is clean. Becomes the new cold layer.
Job-change recovery	A warm contact who moves to a new ICP-fit company gets a personalised welcome ad. Only possible with the owned database, and one of the highest-ROI plays in B2B.

3. The product-market map (the filter on everything)

Short version: one product leads in each market. Anything off this map is waste, full stop.

This was locked with leadership and it filters every recommendation in the doc. We don't run a product in a market where it isn't the agreed play, because off-map spend is waste by definition (the audit found Workforce running in France, and LATAM Workforce buried inside the Spanish account).

MARKET	FLAGSHIP	SECONDARY	OUT OF SCOPE
UK	Flow (hotels first, QSR/FSR second)	Compliance (enterprise hotels, low-volume high-value)	Workforce, Easilys, Inpulse
Iberia (ES + LATAM)	Workforce (enterprise + QSR, strongest local play)	Flow ES, Compliance pilot	Easilys
France	Easilys (contract catering, public sector, healthcare)	Flow (retargeting-led brand build)	Workforce
Nordics	None (no 2026 budget)	Piggyback UK Flow by geo only	Everything else

The competitive picture sharpens the message per market. France is a point-of-sale fight (Lightspeed dominant, Skello present), so Easilys competes there. Iberia competes against HR and payroll tools, so Workforce runs on the HR angle. The UK competes against training and LMS tools, so Flow leads with training-to-execution. Brand defence matters most where rivals bid on our terms, which we've confirmed in France on Bing.

4. The teardown: what we pause and archive

Short version: export everything for the record, then pause and archive the lot, so we build clean.

Before any new campaign goes up, every existing campaign on every account is exported, paused, and archived under a dated label. The history is kept, but nothing old keeps spending or polluting the new structure. The audit found broken conversions still set as primary, polluted audiences, wrong-product campaigns and inherited bid strategies, and the only reliable way to be free of all of it is to stop everything and start clean. Brand defence is the one exception: we reactivate it rather than rebuild it (its history is the baseline we're protecting), then rename it to the convention.

Do this first, on every account: export the current campaign, ad-group, keyword, ad and conversion reports to the dated audit folder. That's the rollback record and the before-state for the eventual re-baseline.

ACCOUNT	WHAT TO DO
Google UK (664-360-5314)	Pause and archive all GROW/CORE campaigns except brand defence. Brand defence reactivates separately. The generic Maximize Clicks campaigns are killed, not rebuilt.
Google ES (656-420-5389)	Pause and archive everything, including the LATAM-nested campaign (off-map for this account) and all 4 Display campaigns.
Google FR (720-427-387)	Pause and archive all 12 Easilys variants and the keyword spread. Keep the 542-keyword Easilys negative list for reuse. The conversion cleanup (18 removed actions still set as primary) happens before the new conversions are trusted.
Microsoft UK (G1457275)	Already fully paused. Archive and rebuild; don't un-pause as-is (wrong target CPA, time zone, currency).
Microsoft FR (G145P17V)	Pause and archive the 42 auto-imported campaigns (includes 4 off-map Workforce ones). Keep the 250-keyword Easilys negative list, currently orphaned.
LinkedIn (EN, ES, FR)	Pause and archive all ABM and legacy ad sets; the narrow Top-100 sets are killed. Keep the FR TOP 100 v2 matched audience (1.28M), the proven lead-gen forms, and the training-gap creative. Decouple the 5-year-old legacy audiences.
Meta (FR live; ES, UK paused)	Pause and archive all campaigns, exclude Audience Network, fix or replace the pixel, delete the 4 zombie pixels. The native lead-form structure and the growth-team video carry into the rebuild.

Archive everything under one searchable label, ARCHIVED-2026-06-PreRebuild (or each platform's nearest equivalent), so the old structure can be found if ever needed and stays clearly separate from the new build.

Before building (leadership, not RevOps)

A terminated contractor still has Super Admin on both Microsoft accounts via a personal Gmail, and there's stale personal-Gmail editor access on the Google accounts. These revocations are leadership actions and should be closed before the rebuild, so the clean account can't be edited by former contractors. Listed in section 14.

5. Google Ads, by market

Short version: four campaigns per market (brand, flagship search, generic, remarketing), with search doing the heavy lifting.

Account settings get fixed first: correct time zone and currency, auto-apply recommendations off (the audit found silent Google-side targeting expansion), and Search Partners and Display unticked on search campaigns.

How the budget splits within Google. Search dominates, because that's where the pipeline is. The automated surfaces stay capped, so Performance Max can't quietly eat the account.

CHANNEL	SHARE	ROLE
Search (flagship + generic)	65-70%	The engine. Pain-point and category keywords by product and region.
Performance Max	10-15%, capped	Kept but constrained, one clean conversion. FR only at launch (it already converts there).
Display retargeting	5-10%	Remarketing only. No cold display (awareness lives in events, trade media, SEO).
Branded search	5-10% cap	Defensive, capped so brand doesn't absorb new-demand budget.

Bidding on competitor brand terms, if we do it at all, goes in a separate, clearly-labelled quarantine campaign with its own small budget and kill criteria. Its economics are usually poor, so we keep them visible and stop them contaminating the campaigns that work. Opt-in per market, and only where a competitor is bidding on us first.

One note on the search campaigns: they send clicks to a landing page, not a native form, so the page has to match the ad and convert. The page builds and any conversion-rate work sit with the website workstream; here we just make sure each campaign points at a page that's fit for purpose before it spends.

Settings for all three Google accounts.

- Time zone: London (UK), Madrid (ES), Paris (FR). It locks after the first campaign, so verify before building.
- Currency: GBP (UK), EUR (ES, FR). Don't mix the UK account into EUR.
- Auto-apply recommendations: off, everywhere. Manual review is policy.
- Conversions: import Lead MQL and Lead SQL from HubSpot, set Lead MQL as the single primary. Remove every legacy conversion still flagged primary before launch.

Google UK (Flow flagship, Compliance secondary)

CAMPAIGN	TYPE / BID	WHAT IT DOES	CONVERSION
UK-HOTEL-Brand-Brand-AlwaysOn-Q226	Search. Manual / impression share (~90%+ on exact brand)	Brand and brand+product terms. Reactivated then renamed. Our cheapest converter historically (€13.64 CPA).	Demo Form Submission
UK-HOTEL-Flow-Search-AlwaysOn-Q226	Search. Maximize Clicks + CPC cap, then Maximize Conversions	Flagship. Ad groups by theme (rota, labour cost, hotel ops), Flow hotel messaging. Fresh negative-keyword list (UK had none).	Lead MQL
UK-QSRFSR-Flow-Search-Generic-Q226	Search. Maximize Clicks + cap	Generic high-intent category terms for QSR/FSR. Tighter budget than flagship.	Lead MQL
UK-HOTEL-Flow-Retarget-AlwaysOn-Q226	Display / Demand Gen remarketing. Small budget	Remarketing to Flow visitors and engaged audiences. The 69,000 unused curated exclusions get attached. Replaces reach-buying.	Lead MQL / Demo

Compliance (enterprise hotels) runs as an ad group inside the flagship, or a small fifth campaign only if budget allows. It's low-volume high-value, so it doesn't need its own always-on campaign at launch.

Google ES (Workforce flagship, Flow ES secondary)

CAMPAIGN	TYPE / BID	WHAT IT DOES	CONVERSION
ES-QSR-Brand-Brand-AlwaysOn-Q226	Search. Manual / impression share	Brand terms. New build (ES brand was never isolated before).	Demo Form Submission
ES-QSR-Workforce-Search-AlwaysOn-Q226	Search. Maximize Clicks + cap, then Maximize Conversions	Flagship. Ad groups: workforce management, labour cost, scheduling. HR-angle messaging. Desktop-weighted (mobile burned €6,593 for zero, so launch with a mobile bid cut). 4 negative lists reattached.	Lead MQL
ES-QSRFSR-Flow-Search-Generic-Q226	Search. Maximize Clicks + cap	Flow ES secondary, lighter feature claims than UK. Generic category demand. Smaller budget.	Lead MQL
ES-QSR-Workforce-Retarget-AlwaysOn-Q226	Display remarketing. Small budget	Remarketing to Workforce visitors. Exclusion list attached (ES had none despite Display running 666-782K impressions).	Lead MQL

LATAM has no separate Google account; it sits under ES. Run it as its own Workforce search campaign (LATAM-QSR-Workforce-Search-AlwaysOn-Q226) with its own geo-targeting and a small budget (€0.6-1K/month), not buried inside ES campaigns like before.

Google FR (Easilys flagship, Flow secondary)

CAMPAIGN	TYPE / BID	WHAT IT DOES	CONVERSION
FR-CC-Brand-Brand-AlwaysOn-Q226	Search. Manual / impression share	Brand and Easilys-brand terms. Genuinely incremental here: rivals bid on our brand on Bing.	Demo Form Submission
FR-CC-Easilys-Search-AlwaysOn-Q226	Search. Maximize Clicks + cap, then Maximize Conversions	Flagship, consolidated from 12 variants into one. Ad groups by theme, geos as ad groups (not separate campaigns). 542-keyword negative list reattached.	Lead MQL
FR-CC-Easilys-Demand-AlwaysOn-Q226	Performance Max (single, clean)	One PMax on a single clean conversion. PMax converted in the audit whenever the signal reached it. Launch paused until the conversion cleanup is done, then enable.	Lead MQL
FR-CC-Easilys-Retarget-AlwaysOn-Q226	Display remarketing. Small budget	Remarketing plus the Flow retargeting-led brand build (Flow is FR secondary). FR exclusion lists attached.	Lead MQL

TRACKING PREREQUISITE, FR CONVERSIONS (OURS). The FR conversion layer needs a cleanup in the Google Ads UI (29 actions, 18 removed, 4 still flagged primary, Google warns the goal can't optimise). This is ours to do and must finish before FR switches to conversion bidding. Until then, FR runs on Maximize Clicks with caps.

6. LinkedIn Ads, by market

Short version: three campaigns per market, broad audiences, native lead forms. Broad wins; narrow lists don't.

This is the lesson the audit proved most clearly: broad, correctly-targeted audiences win, and narrow hand-built lists fail. Account-based work is rebuilt as three tiers instead of narrow lists, with every always-on tier above the 50,000 floor. Lead Gen Forms are the default capture format, because the audit showed native forms beat landing pages, and because the form makes the campaign name the attribution.

One correction from the live rebuild: the Insight Tags fire on all three accounts, but LinkedIn is not actually connected to HubSpot yet (no CRM sync, and Lead MQL/SQL aren't attached at ad-set level). So part of the work here is the connection, not just the campaign structure. It's in section 14 (items 5, 10, 12).

The format mix within the LinkedIn budget. LinkedIn isn't only a lead-form harvester. The mix balances capture with the demand-creation formats the project proved, all inside the existing 70/30 evergreen/growth split (a reweighting, not new budget).

FORMAT	SHARE	ROLE
Sponsored Content + Lead Gen Forms	40-45%	Primary capture. Whitepaper and product ads to native forms.
Conversation Ads	30-35%	Cheap, converts well. First two lines read as relevance, not a pitch. Strong for warm and MQL audiences.
Document Ads	15-20%	The French whitepaper-to-engagement pattern, now the template across markets. Feeds the retargeting pool.
ABM	10% max	The three tiers below. Justified on pipeline or cut.

The three campaigns per market.

CAMPAIGN	FORMAT / BID	AUDIENCE	CONVERSION
{REGION}-{ICP}-{PRODUCT}-LeadGen-AlwaysOn-Q226	Sponsored Content, native form. Maximize leads, cost cap later	Broad ICP, above 50K, by title + industry + size. The workhorse. Training-gap creative (4-ad matrix) carries forward.	Lead MQL
{REGION}-{ICP}-{PRODUCT}-Demand-Sponsored-Q226	Sponsored Content (image / document). Lowest cost	Mid-funnel, broad ICP. Builds the retargeting pool. Content-led, not gated.	Lead MQL / engagement
{REGION}-{ICP}-{PRODUCT}-Retarget-AlwaysOn-Q226	Sponsored Content retargeting. Cost cap	Engaged audiences and site visitors. Stays present between nurture emails.	Lead MQL / Demo

How ABM is rebuilt (three tiers, inside the structure above).

TIER	AUDIENCE	APPROACH
Tier 1: Named accounts	10-30 priority logos (sales-chosen)	One-to-one Sponsored Content + retargeting. Higher CPC accepted, deal size justifies it. Measured on meetings booked. Sales sequences run in parallel (not paid).
Tier 2: Best-fit (the workhorse)	50K+ best-fit by title, industry, size	The proven broad model: native forms + sponsored content. Measured on cost per lead through to pipeline. Where budget and pipeline concentrate.
Tier 3: Broad ICP + lookalikes	Widest correct ICP plus lookalikes of converters	Lowest cost, feeds the funnel and the retargeting pool, graduates contacts into Tier 2. Lookalikes from closed-won once tracking feeds revenue back.

The full ABM build detail (the narrow-versus-broad proof, the build standards every ABM campaign must meet, and worked tier examples for French Easilys) is in Appendix A.

Per-market notes.

- **FR (Easilys):** build on the existing FR TOP 100 v2 matched audience (1,285,080 members, >90% match) that sat unused while a narrow wrong-vertical list ran. Filter it for Tier 2, use it broad for Tier 3. Use the proven Efficca Ope v2 form (21 leads) as the template. Decouple the 5-year-old legacy audiences.
- **UK (Flow):** the training-to-execution creative is sound and carries forward. The audit found clicks producing zero leads, which points to a landing-page failure, so default to native forms and check the form fires before spend. Reuse the gamified-learning form as the EN reference.
- **ES (Workforce):** reuse the two proven ES forms (whitepaper and MX, 5 leads each) on a broad Workforce audience. The ES ABM set was paused with little data; don't revive the narrow approach.

TRACKING PREREQUISITE, LINKEDIN (MOSTLY OURS). Attaching Lead MQL/SQL at ad-set level, reactivating the FR event, and clearing the dead events from the account pools are ours, in LinkedIn Campaign Manager. The one genuine RevOps piece is connecting and confirming the HubSpot to LinkedIn integration so lead and conversion events sync. Section 14.

7. Microsoft Ads, by market

Short version: treat Microsoft as a new channel, not a fix. Two campaigns per market: brand and generic.

Microsoft was near-zero spend a year ago, so we build it properly rather than patch it. The UK account is rebuilt from corrected settings (never un-paused as-is), France is rebuilt around Easilys and brand after the auto-imported sprawl is archived, and Spain gets a new account in Q3.

Settings.

- UK (G1457275): time zone Helsinki to London (may need a support ticket, often locked), currency EUR to GBP, remove the unrealistic €6 target CPA, build a negative-keyword list (account had none).
- FR (G145P17V): archive the 42 auto-imported campaigns, attach the orphaned 250-keyword Easilys negative list, remove the ES audience provisioned here by mistake.

CAMPAIGN	TYPE / BID	WHAT IT DOES	CONVERSION
UK-HOTEL-Brand-Brand-AlwaysOn-Q226	Search. Manual / impression share	Brand defence, reactivating the paused UK presence. Catches the 5-10% of branded search Google misses. Add negatives for the competitor terms we were wrongly bidding on.	Demo Form Submission
UK-HOTEL-Flow-Search-Generic-Q226	Search. Maximize Clicks + cap	Generic high-intent, Microsoft-specific keywords (not the generic terms that burned €1,570 for zero). Flow flagship.	Lead MQL
FR-CC-Brand-Brand-AlwaysOn-Q226	Search. Manual / impression share	Brand defence. Critical: Skello, Sage and Eurecia bid on our brand on Bing (impression share to 56%). Add negatives for the Workforce-mention queries eating FR brand budget.	Demo Form Submission
FR-CC-Easilys-Search-Generic-Q226	Search. Maximize Clicks + cap	High-intent Easilys. The FR Easilys brand term already converts at ~€143, the one positive signal, so build around it. 250-keyword negative list attached. No Audience Network.	Lead MQL

Spain Microsoft (Q3): no ES account exists today and there's uncovered ES demand on Bing. Open one in Q3, funded by reallocating ES Workforce Google waste, mirroring the two-campaign structure. Not part of the launch build, just on the roadmap.

8. Meta Ads, by market

Short version: three campaigns per active market, all native lead forms, all pointed at leads, never traffic.

Native forms were three times cheaper than landing pages in the audit. Everything points at leads, never traffic or awareness (which was 73% of the wasted spend). France is the live base; UK and ES are rebuilt as they come online. The pixel gets fixed and the four zombie pixels deleted before any scaling, and Audience Network is excluded everywhere (it leaked spend on both Meta and Microsoft).

CAMPAIGN	OBJECTIVE / BID	AUDIENCE	CONVERSION
{REGION}-{ICP}-{PRODUCT}-LeadGen-Native-Q226	Leads, native Instant Form. Lowest cost, cost cap later	Broad ICP. The €26 vs €84 winner. Audience Network excluded. Form name is the attribution.	Lead (then Lead MQL once traced)
{REGION}-{ICP}-{PRODUCT}-Retarget-Conversion-Q226	Conversion remarketing. Cost cap	Site and engagement audiences. Replaces reach-buying. Needs the working pixel.	Lead / Demo
{REGION}-{ICP}-{PRODUCT}-Demand-Video-Q226	Demand Gen / video, inside the lead-form structure	The growth-team video runs here, where it converts, not as standalone awareness. One clean audience.	Lead

France launches first. UK and ES come online once their pixels are confirmed firing and their forms built. Scaling any market waits on tracing the native-form leads through to pipeline, so the format is proven on value, not just cost.

TRACKING PREREQUISITE, META PIXEL (OURS). The main pixel was dead ~27 days and four others are inactive; lead forms only survive by bypassing the pixel. Restoring the main pixel, deleting the four zombie pixels, and confirming Lead and Lead_SQL fire are ours, in Events Manager, before remarketing or conversion bidding. Section 14.

9. The nurture engine paid feeds into

Short version: paid brings the lead in, nurture turns it into a deal. Today most leads go cold. This fixes that.

Paid is half the engine. The other half is what happens after a lead comes in. For a slow, considered purchase like ours, this is where the deal is actually won or lost, and right now it's the weak point: most non-demo leads dead-end. The detailed build lives in the separate nurture-to-SDR workflow doc. The design, the model and the targets are here.

Lead > MQL > SAL > SQL > Opportunity > Customer

The lifecycle stages (defined the same way for marketing and sales). Every contact sits in one stage, defined the same way for both teams, so "MQL" means the same thing to marketing and to sales. When stages are loosely defined, leads dead-end and the reporting lies (the audit showed exactly that). So the shared definition comes first; everything below depends on it.

STAGE	WHAT IT MEANS	OWNER	WHAT MOVES IT FORWARD
Lead	Known contact, not yet qualified	Marketing	Fit + behaviour scoring
MQL	Hits the score; ready for a human	Marketing	Score crosses the threshold (set by scoring, never by hand)
SAL	Sales has accepted it as worth working	SDR	SDR accept; a reject sends it back to nurture with a reason
SQL	A real opportunity (need, fit, timing)	SDR / Sales	Discovery confirms it
Opportunity	Open deal in the pipeline	Sales	Stage progression
Customer	Closed won	Sales	Signed

How a lead becomes an MQL (fit + behaviour). A lead qualifies on two scores together, not one action. Fit asks "is this the right company and person." Behaviour asks "are they actually interested." The interim model runs now on HubSpot scoring and swaps to the RevOps model later, with no change to the design.

SCORE	EXAMPLES	THE PRINCIPLE
Fit (who they are)	Title / persona, company size 20+, industry on the map, region	Hard-gated: sub-20-employee (matching the inbound qualification gate) or off-map caps the score, so juniors and wrong-vertical leads can't reach MQL on behaviour alone.
Behaviour (what they do)	Pricing-page visit, 2+ downloads, demo interaction, email engagement, repeat visits	High-intent actions (pricing, demo) count far more than a single email open.

The threshold is set with sales and SDR input, not by me alone, because they own what "ready" means downstream. We review it against the actual MQL-to-SQL rate and tune it, so the bar tracks real quality. And MQL is set by scoring, never flipped by hand, because manual delay misses the platform match window that value-based bidding depends on (section 14, item 11).

Routing, and the handoff to sales.

- **Tag on entry:** every paid lead is flagged with its campaign and a paid-lead marker, so it's visible for routing and reporting. (This is why naming and UTMs matter so much.)
- **No more dead ends, now live:** a central routing workflow creates, qualifies and routes the non-demo leads (whitepaper, guide, event, product-page) to an owner, splits current customers to Customer Success, dedups against open leads, and blocks the clear non-fits. The residual gap is that only demo requests alert an SDR today, so non-demo leads can sit unworked; closing that, a queue view or daily digest, sits with RevOps and the SDR side (section 14, item 8).
- **Graduate to sales:** once automatic scoring is live (RevOps Phase 2), a sales-ready lead creates a high-priority, paid-labelled task in the SDR queue, with a same-day response as the target. Since SDRs are at 63% of target, it stays priority and a soft target, not a punitive timer. Named accounts can route to a LinkedIn-connect task instead.
- **Accept or reject:** the SDR accepts (it becomes SAL) or rejects with a reason. Rejected leads go back to nurture with that reason, which is also the cleanest lead-quality signal back to paid.
- **Remarketing keeps it warm:** paid remarketing stays present between nurture emails. The audience sync waits on the tracking fix; the nurture

design doesn't.

Three layers of personalisation. Every sequence and every ad is personalised on three things, because a generic "Mapal" message converts worse than a specific one:

- **Product:** the module they showed interest in.
- **Market:** UK, Iberia, France, LATAM, with the language and local customer proof.
- **ICP:** hotels vs QSR/FSR vs contract catering vs healthcare.

So someone who downloaded a Workforce guide in Spain gets Workforce nurture, in Spanish, with a Spanish customer story. Not generic content.

Which nurture a lead enters, and what speeds it up. Nurture isn't one sequence. It branches by the module the lead cared about and by their role, because our products speak to different buyers.

TRACK	MAIN BUYER	WHAT WE TALK ABOUT
Flow / Compliance	HR, Operations	Onboarding speed, training consistency, compliance gaps, cost of staff turnover.
Workforce	Operations Director, Finance	Labour cost, scheduling, AI-driven forecasting (HR angle in Iberia).
Easily	Finance, public-sector and contract-catering leads	Food-cost control, kitchen and procurement efficiency, public-sector case studies.

Inside each track, a four-week sequence: week 1 a relevant customer story (no pitch), week 2 content that helps build the internal business case, week 3 proof from a similar business in their market and language, week 4 a low-friction demo offer. Non-converters drop to a monthly long-term nurture rather than going cold.

On top of the sequence, behaviour triggers an accelerated response:

SIGNAL	WHAT HAPPENS
Visits the pricing page	Immediate case-study email with a demo offer, plus an ROI retargeting ad.
Downloads 2+ pieces	SDR alert for warm outreach, plus a case-study ad within 48 hours.
Opens 3 emails, no clicks	Switch format (try a video or a different CTA). They're reading but not engaging.
Demo no-show	Conversation Ad within 7 days to re-open it.
Goes quiet for 45 days	Re-engagement sequence keyed to their original problem.

Who we leave out of nurture. Same discipline as paid. We suppress a contact from active nurture when they're an open opportunity, a current customer (they belong in cross-sell, a different motion), unsubscribed, or already in a live SDR conversation. So we never email someone who's already being worked or who has opted out. Re-engagement catches non-converters rather than letting them churn out of the database.

Retargeting wired to behaviour, not just "visited the site." Remarketing here follows what the lead does, so paid reinforces exactly what nurture is doing at that moment: the whitepaper downloader sees a case-study ad within 48 hours, the demo no-show sees a Conversation Ad within 7 days, the pricing-page visitor sees an ROI ad, the high-health customer sees a cross-sell ad. It's gated behind the working pixel and tag. The useful part for the build: the conversion a campaign points at (Lead MQL) is the same signal that triggers the routing and the retargeting, so building on the right conversion now means all of it has clean inputs the moment it switches on.

How we judge nurture. On pipeline, not activity. In order of importance:

METRIC	WHAT IT PROVES	TARGET
SAL-to-booking rate	Quality of what we hand sales	From sub-15% toward 20% by end of 2026
MQL-to-SQL, nurtured vs not	The lift nurture actually creates (the core proof)	Nurtured clearly above non-nurtured
Time-to-SQL	Whether nurture shortens the cycle	Down over time
Pipeline influenced	Revenue the programme touches	Up, reported in HubSpot
Engagement and unsubscribe	Programme health	Engagement up, unsubscribe under 0.3% per send

We don't report sends, opens on their own, or "leads touched." Those are activity, not outcome. The programme is judged on pipeline and on the nurtured-versus-non-nurtured lift, the same revenue-first bar the paid side is held to.

How we build it (Q3 to 2027). The nurture build ships value early and scales in stages, the same as the paid rebuild. It never tries to finish at stage one.

PHASE	WHEN	WHAT SHIPS
Quick wins	Q3	Audit and fix current workflows, kill broken ones, fix UTM and source attribution, build the scoring model with sales and SDR input, ship the 2-3 highest-value sequences (post-demo no-show, top whitepaper follow-up).
MVP 1	Q4	Full track-based architecture live, the product / market / ICP personalisation in place, dashboard reporting going.
MVP 2	2027	Progressive profiling, predictive scoring, multi-touch orchestration, account-level nurture for ABM.

How nurture is run and requested.

- **Ownership:** the strategy and execution sit with paid and lifecycle. Growth marketing and product marketing bring content, ICP intel and market input. Sequence design, triggers, frequency and orchestration are decided centrally so the system stays coherent.
- **Requests through intake, not chat:** nurture and content requests come through the same intake form as paid campaigns (scoped for nurture), so each request carries what's needed to build without back-and-forth.
- **One taxonomy:** a one-off email blast either fits the nurture taxonomy or runs as a clearly separate broadcast. It doesn't get bolted onto a live sequence.

10. The expansion engine: cross-sell and up-sell

Short version: acquisition wins a logo, expansion grows it. In a multi-module business like ours, the cheapest pipeline we have is the customers we already own. This is the motion the acquisition rules keep pointing to when they exclude customers.

Two of the loops are already in this plan: paid captures the lead, nurture turns it into a deal. The third loop is what happens after the first sale. A customer who adds a second module and rolls out across their estate is worth several times their first contract, converts far better than a cold prospect, and costs a fraction to reach because we already own the relationship and the data. Across the plan we deliberately suppress current customers from acquisition campaigns and acquisition nurture, with the note that they belong in cross-sell, a different motion. This section is that motion.

Customer > Adopt (uses what they bought) > Cross-sell / up-sell > Expanded account

The expansion loop. It only runs on healthy, adopting accounts; an unhappy customer is a retention problem, not an up-sell target.

Two motions, defined. Cross-sell is a customer on one module taking an adjacent one (a new product line into an existing account). Up-sell is growing what they already have: more sites, more seats, a higher tier, add-ons, or a longer term. Both are measured as expansion revenue and net revenue retention, not new-logo pipeline.

Who owns what. The expansion sale is owned by Customer Success and account management (and sales on the largest accounts). Paid and lifecycle own the demand layer: surfacing the right opportunity at the right moment, warming the account, and handing a qualified expansion signal to the account owner. We generate and shape expansion demand; we don't close it. That keeps us out of CS's lane while making their job easier.

The cross-sell map (who is a natural candidate for what). It follows the same product-market map as acquisition, and the same product / market / ICP personalisation, so a Spanish Workforce customer gets a Spanish Compliance cross-sell with Spanish proof.

CURRENT PRODUCT (ENTRY)	NATURAL NEXT MODULE	THE ANGLE
Flow (UK hotels, the largest base)	Compliance, then Engage	"You already train your teams consistently. Here's how to close the compliance and audit gap at the same time." Flow and Compliance are often bought together.
Workforce (Iberia)	Compliance + Engage	"You've fixed scheduling and labour cost. Here's the compliance and engagement layer on top of the same teams."
Easilys (France, contract catering)	Flow	"Your kitchens and procurement run on Easilys. Here's how to bring training and onboarding into the same operation."
Multi-module / Key Accounts	MapalOS consolidation	The platform story: more modules on one operating system, one login, one data layer. The largest-account expansion play (the LATAM key-account upsell is the reference proof).

The up-sell paths. Site expansion is the big one in multi-site hospitality: a group live in some sites rolls out to the rest of the estate. Then seat or user expansion within a site, a tier upgrade or add-on, and the renewal uplift. Renewal is the natural moment for all of these.

The trigger logic (the heart of it). Expansion fires on a combination of account health, product usage and timing, never on tenure alone.

TRIGGER	WHAT HAPPENS
Health gate (the one hard rule)	Only green-health accounts (the CS health score) enter cross-sell or up-sell. Amber and red go to adoption or save motions, never an ask for more money. This rule sits above every trigger below.
High adoption + tenure on the owned module	Enters the cross-sell sequence for the adjacent module (e.g. Flow live 6 months, high usage, green health, no Compliance, then the Flow-to-Compliance sequence).
Low adoption of an owned module	Adoption nurture first, not cross-sell. You can't expand a customer who isn't using what they already bought; protecting the base comes before growing it.
Nearing a seat or site limit	Up-sell sequence (more seats, estate roll-out), with an account-owner alert.
Renewal window (90 / 60 / 30 days)	Value-reinforcement sequence that surfaces the expansion opportunity alongside the renewal.
New senior stakeholder at the account	A new decision-maker appears at a customer (org change), introduce the adjacent module to them. (This is the customer-side twin of the job-change recovery play in section 2.)

The nurture layer (lifecycle-owned). Customers run in their own audiences and sequences, completely separate from prospect nurture. A customer-only content stream stays always-on at low pressure: product updates, peer benchmarks, ROI and value content, roadmap previews, and customer-summit and event invites. On top of that sit the gated cross-sell sequences (current product to adjacent module, value-led not feature-led, four to six touches, ending in a soft "worth a conversation with your account manager?" that creates a CS task, not a demo form), the renewal 90/60/30 sequence, and the adoption nurture that keeps the base healthy enough to expand.

The paid layer (this is the cross-sell into PPC). Paid becomes a multiplier on the customer database, the same database-activation mechanism described in section 2, pointed at customers instead of prospects.

- **Customer audiences:** built from HubSpot customer segments by module owned and health score, pushed to LinkedIn Matched Audiences, Google Customer Match and Meta Custom Audiences. Each cross-sell campaign suppresses the module the account already owns, plus amber/red accounts and any account already in an open expansion opportunity.
- **Cross-sell creative:** module-adjacent and proof-led ("you run Flow, here's what Compliance adds"), carrying peer and logo proof from similar accounts.
- **Up-sell / estate roll-out** ads to multi-site groups, on the "live in some sites, here's the rest of the estate" message.
- **Behaviour-triggered:** a green-health customer who visits an adjacent-module page sees a cross-sell ad. This is the customer half of the behaviour-triggered retargeting already in section 9, now made explicit.

The whole paid layer depends on the customer-segment sync from HubSpot to the platforms and on product-usage and health data being available to build the audiences. We design it now and switch it on behind that data, exactly as we do for prospect database activation.

The handoff. When a green-health account crosses the expansion threshold, it creates a high-priority, paid-labelled task for the account owner (CS or account management), carrying the suggested module and the signal that fired it. This mirrors the nurture-to-SDR engine in section 9, but routes

to whoever owns the account rather than to an SDR, and it is never a raw demo form.

How we judge it. On expansion revenue, the same value-over-volume and lift-not-last-click discipline as the rest of the plan. A cross-sell holdout is even more important here than on cold acquisition, because these customers might expand anyway: hold a matched set of green-health accounts out of the cross-sell ads over a fixed window and compare expansion against the accounts that saw them, so we only credit paid with the lift it actually creates.

METRIC	WHAT IT PROVES
Expansion revenue and net revenue retention	The headline, owned jointly with CS
Cross-sell rate (share of customers on 2+ modules) and movement over time	Whether the multi-module motion is working
Up-sell: estate roll-out rate and average uplift at renewal	Within-account growth
Marketing-influenced expansion opportunities created	What the customer nurture and paid layer contributes
Adoption health of owned modules	The leading indicator; non-adopters can't be expanded

When we build it. This is post-launch and depends on product-usage and health data plus the customer-segment sync, so it is a Q3 to 2027 build, not part of the paid launch. Quick wins in Q3: stand up the customer-only audiences, the two highest-value cross-sell sequences gated on green health (Flow-to-Compliance and Workforce-to-Compliance), and the renewal 90/60/30 sequence, with matching customer Matched Audiences. Q4: the full cross-sell map live, the up-sell and estate-roll-out triggers, and paid customer-match cross-sell across markets. 2027: predictive expansion scoring and account-level orchestration, with the ABM Tier 1 treatment (Appendix A) pointed at key-account expansion.

DATA PREREQUISITES (BEYOND THE ACQUISITION FIX-LIST). Expansion adds two dependencies to the section 14 items, now tracked there as owned items 14 and 15: product-usage data flowing into the CRM (owner: RevOps, target Q3) so adoption and ceiling triggers can fire, and the customer-health score exposed to marketing (owner: Customer Success, target Q3) so the green-health gate and suppression can be enforced. Neither blocks the acquisition rebuild.

11. Creative and testing

Short version: a few strong, tested angles per product beats a pile of one-off ads. Once structure is right, creative moves the cost per lead.

The naming convention has a creative slot ({Format}-{Angle}-V##) for exactly this reason: every asset is traceable to an angle and a test. Ads lead with the reader's problem, then how we solve it, then proof, not generic product claims.

The angle per product.

PRODUCT	LEAD ANGLE	PROOF
Flow (UK, hotels)	Training that turns into consistent on-shift behaviour; the cost of staff turnover	The training-gap creative (4-ad matrix) and the gamified-learning asset; hotel stories
Workforce (Iberia)	Labour cost and scheduling on the HR angle (we compete with HR and payroll tools, not POS); AI forecasting	Whitepaper and the proven ES forms; same-size local proof
Easily (France)	Food-cost control, kitchen and procurement efficiency; public-sector cases	The document/whitepaper pattern that worked in France; public-sector references
Compliance (UK secondary)	Compliance-gap closure and audit-readiness for enterprise hotels	Enterprise hotel proof; low-volume, high-value framing

What each format is for.

FORMAT	ROLE
Single image + native form	Primary capture. Problem-led hook to a low-friction form.
Document / whitepaper	Mid-funnel education that builds the business case and feeds retargeting.
Conversation / messaging	Warm and MQL re-engagement. First two lines read as relevance, never a pitch.
Video	Demand and view-through inside a lead-form structure, not standalone awareness.

How we test.

- **One thing at a time:** test the hook, the image, or the CTA, never all three at once, so the winner is clear. The V## in the name tracks the version.
- **Judge on the funnel, not the click:** CTR is an early signal, but the decision metric is cost per lead through to pipeline. A high-CTR ad that brings junk is a loss.
- **Refresh on fatigue:** when frequency climbs and CTR drops on a still-live audience, rotate the angle rather than spend more. Fatigue is usually the cause, not budget.
- **Sourced through intake:** creative and copy come from growth marketing and design via the intake form. The angle and format call sits with paid, the production with the content owners.

12. How we measure success

Short version: we judge spend on pipeline, not platform-reported clicks. HubSpot pipeline is the truth.

Platform dashboards optimise to their own conversions, which are directional at best. The real scoreboard is HubSpot pipeline, closed-loop through the offline conversion upload.

What every campaign is held to. These are planning ranges to firm up once we have a reliable cost per lead by channel, not promises. The point is that every campaign carries them from launch (the interim ceilings in section 2 are the floor).

Metric	Role	Reference
Cost per lead, by channel and market	The launch-phase efficiency gate	Native social forms ran €26-55 in the audit; ceilings €60-90 social, €120-150 search
MQL volume	Demand we generate	By channel and market
MQL-to-SQL rate	Lead quality (matters more than volume)	What value-based bidding optimises toward
Cost per SQL	True acquisition efficiency	Replaces cost-per-click as the headline number
Pipeline and win rate	Revenue contribution	Paid's trailing-12-month win rate is 15.8%, the baseline to hold or beat as we scale

Value over volume. Because this is a slow, considered purchase, the account optimises for lead value, not lead count, by pushing values back from HubSpot (MQL €500, SQL €2,000, Opportunity €10,000, Closed Won €25,000). Volume bidding rewards cheap, low-intent leads; value bidding shifts delivery toward the people who become real deals. It's the default for the flagship and lead-gen campaigns (mechanism in section 2; prerequisites in section 14).

Measure lift, not last-click. A channel is judged on the lift it actually creates, not the conversions a platform claims. The cheapest last-click cost per lead can hide spend that would have converted anyway (branded search is the classic case), so we build incrementality into the plan instead of bolting it on later.

Test	When	What it answers
Brand-defence geo holdout	At launch	Is brand-defence spend incremental, or just harvesting demand that would convert anyway? Hold brand out in a set of matched regions over a fixed window (a few weeks) and compare against the regions where it keeps running. A well-known version of this test found eight-figure "working" brand spend was largely not incremental, so the holdout pays for itself.
Meta value holdout	Q3, after leads are traced to pipeline	Does Meta lead spend drive incremental pipeline, measured on value not lead count? Earned as volume grows, not front-loaded onto a rebuilding account.

Move budget to what works.

- **Reallocate, don't use-it-or-lose-it:** budget moves monthly toward the channel, market and campaign producing pipeline-qualified leads at the best cost per SQL, rather than being spent because it's there. The social growth shares are ceilings that revert to evergreen if unused.
- **Read year-on-year, not just month-on-month:** hospitality demand is seasonal, so we compare to the same period last year before reading a change as real.
- **Two scenarios, not one forecast:** the redeployment is planned as a conservative and an optimistic case, consistent with the report, so we're judged against a range.

How often we look.

Cadence	What
Twice weekly (interim)	Spend, clicks, leads, CPL by campaign. The instrument for the riskiest weeks while bidding runs blind.
Weekly	Hygiene across the ~33 campaigns: pacing, search terms, negatives, anything odd.
Monthly	Strategic review with the lead-quality report and sales feedback: which campaigns produced accepted leads, MQL-to-SQL by source, where budget moves next.

13. Budget and campaign count

Short version: no new budget. We fund the rebuild by redeploying the waste, into about 33 campaigns one person can actually run.

The numbers below are planning ranges, to firm up once we have a reliable cost per lead by channel. Search stays fully evergreen; the social platforms split evergreen and a ring-fenced growth share (a ceiling, not a floor: unused, it reverts to evergreen).

Platform	UK	Iberia	France	LATAM	Monthly
Google	€4-6K	€4-6K	€4-6K	€0.6-1K	~€14-19K
LinkedIn (70/30)	€3-4K	€3-4K	€1.5-2K	-	~€7.5-10K
Microsoft	€1-2K	Q3 build	€0.8-1.6K	incl.	~€2-3.5K
Meta (75/25)	€0.5-1K	€0.5-1K	€1-1.5K	-	~€2-3.5K

Platform	Per market	Markets	Approx total
Google	4 (brand, flagship, generic, remarketing)	UK, ES, FR (+ LATAM search)	~13
LinkedIn	3 (lead-gen, demand, retarget)	UK, ES, FR	~9
Microsoft	2 (brand, generic)	UK, FR (ES in Q3)	~4
Meta	3 (lead-gen, retarget, video)	FR live, UK + ES as they come online	~7
Total			~33

One person can run weekly hygiene across 33 focused campaigns. The current sprawl can't be run well by anyone, which is the practical argument for consolidating.

14. Tracking and routing remediation, and who owns each

Short version: most of this is ours to execute (the platforms and the tag container); a short set genuinely needs RevOps; the lead-routing engine is already live.

The conversion outage is a website dataLayer break, not an attribution fault, so the bulk of the work sits with us and the website-migration lead rather than RevOps. The HubSpot attribution design is sound. Each row names what is wrong, what is needed, and who owns it, written so the genuine cross-team items can be raised as their own tickets.

#	ITEM	WHAT'S WRONG	WHAT'S NEEDED	OWNER
1	Conversion signal / site dataLayer	The migration moved form conversions onto a HubSpot form event, then renamed the form-id dataLayer keys, so the triggers stopped matching. This is the critical path. The Google Ads conversions are live but untriggered.	Verify the migrated site pushes the form event with the current keys (a live dataLayer test), then realign the triggers in the tag container.	Us + website lead
2	FR Google conversions cleanup	29 actions, 18 removed, 4 still flagged primary; Google warns the goal can't optimise.	Unmark removed actions as primary, remove from account goals, leave only live Lead MQL and Demo.	Us
3	ES Google "Compra"	Removed, yet still primary and in account goals; bidding may chase a dead event.	Unmark from primary, remove from account goals.	Us
4	UK Google "Demo Form Submission"	In "needs attention" but recording 20 conversions; the count isn't trustworthy.	Check the tag, confirm it fires, clear the state so the count can be trusted.	Us
5	LinkedIn conversions + HubSpot connection	Lead MQL/SQL not attached at ad-set level; FR Lead MQL inactive; pools polluted with dead events; the HubSpot to LinkedIn sync isn't connected.	Attach the events, reactivate FR, clear the dead events (ours). Connect and confirm the HubSpot to LinkedIn integration (RevOps).	Us + RevOps
6	Microsoft events	"No recent conversions" on both; the UET tag is healthy but events aren't arriving (same dataLayer root cause as item 1).	UET and config are ours; confirming events flow from HubSpot into Microsoft is RevOps.	Us + RevOps
7	Meta pixel + zombies	Main pixel dead ~27 days, four others inactive; forms survive by bypassing the pixel.	Restore the main pixel, confirm Lead and Lead_SQL fire, delete the four zombie pixels.	Us
8	Non-demo lead routing	Now live: non-demo leads are created, qualified and routed to an owner, customers split to CS, duplicates and non-fits filtered. Residual: only demo requests alert an SDR, so routed non-demo leads can sit.	Routing shipped. Add a working signal on non-demo leads (a queue view or daily digest) and a clean paid-source label on entry.	RevOps (shipped) + SDR side
9	Cross-channel UTM and deal attribution	UTM depth onto deal records unconfirmed; campaign write-back and the Dynamics source-overwrite are the genuine upstream gap.	Confirm UTMs reach deal records and campaign writes back, and stop Dynamics overwriting source, so the naming rolls up as designed.	RevOps
10	Offline conversion upload (LinkedIn + Google)	Value-based bidding can't run without it; the platforms can't learn which leads became deals.	Enable the upload on the platforms (ours); confirm the HubSpot connection that feeds deal-stage values back (RevOps).	Shared, RevOps-gated
11	Automatic MQL / lead scoring	Manual MQL misses the platform match window value bidding depends on. Scoring on fit and intent is RevOps' roadmapped Phase 2.	Make MQL scored or workflow-driven so it fires automatically. Sales signs off the threshold.	RevOps, Phase 2
12	li_fat_id + UTM hidden fields	Without the LinkedIn click ID on the form, offline conversions can't tie back. Most-skipped step, silently kills attribution.	Add li_fat_id and utm_source/medium/campaign as hidden fields on every form. No form ships without it.	Us
13	Attribution windows	Senior buyers don't convert same-session; short windows undercount paid.	Set a 30-day click-through window in each platform (ours); align the HubSpot model window with conservative view-through (RevOps).	Us + RevOps
14	Product-usage data synced into the CRM	Expansion adoption and ceiling triggers can't fire without it. Gates cross-sell, not the acquisition rebuild.	Sync product-usage data into the CRM. Target: Q3.	RevOps / data
15	Customer-health score exposed to marketing	The green-health gate and customer suppression can't be enforced without the health score in marketing's reach. Gates cross-sell.	Expose the customer-health score to marketing. Target: Q3.	Customer Success

Items 1 to 7 are the conversion-tracking work, ours and the website lead's. Items 8 and 11 are the routing engine and its Phase 2 scoring, RevOps owned (8 is live). Items 9, 10 and 13 are the cross-system and optimisation layer. Items 14 and 15 are expansion dependencies for section 10, targeted Q3, and gate cross-sell only, not the acquisition rebuild.

Governance (leadership and marketing-admin, not RevOps)

Close the departed-contractor access across the stack as handover completes. The tag container has personal-Gmail accounts with admin rights and several external-agency publishers (a departed contractor pushed live changes as recently as mid-May), and there is a Super Admin on both Microsoft accounts plus stale personal-Gmail editor access on the three Google accounts. Revoke these, and turn on "require sign-in with work account" where the platform supports it, before the rebuild.

15. The launch order

Short version: governance and teardown first, clean settings second, capture-ready campaigns third, the conversion-bidding switch only once the signal is real.

Lead capture is never blocked: native forms work from day one. Only conversion-based bidding and the remarketing sync wait on the tracking fix.

STAGE	WHAT HAPPENS	WAITS ON
0. Governance + export	Close the access revocations. Export every account to the dated audit folder.	Leadership for revocations; nothing for export.
1. Teardown	Pause and archive everything under the dated label. Keep the named audiences, negative lists and proven creative.	Export done.
2. Settings	Fix time zones, currencies, target CPAs; auto-apply off; import and set the primary conversion; attach negative-keyword and exclusion lists.	Teardown done.
3. Brand + flagship	Reactivate and rename brand defence; build the flagship search per market on interim bidding; build the LinkedIn lead-gen and Meta native-form campaigns (these capture now).	Settings done.
4. Generic + remarketing	Build the generic and remarketing campaigns. Remarketing launches small; full behaviour waits on the conversion-signal fix.	Flagship live.
5. Switch bidding (T+7d)	Once we and the website lead confirm the conversion signal is firing (the dataLayer contract realigned and verified by a live test) and the in-platform conversions are clean, verify events fire for 7 days, then switch flagship and generic from Maximize Clicks to Maximize Conversions, set realistic target CPAs, enable FR Performance Max.	Tracking items 1-7 verified.
6. Sync + scale	Sync HubSpot segments to platform audiences, build revenue lookalikes, scale the native-form model to UK and ES once leads are traced to pipeline.	Tracking healthy; leads traced.

Q3 roadmap (depends on other teams, not part of the launch). Three bigger plays are designed but deliberately held for Q3, because they depend on things outside paid:

- **Paid amplifies SEO.** Once SEO is producing ranking problem-aware content, paid social amplifies the best pieces (paid serves SEO, not the reverse), and SEO becomes the cold top-of-funnel engine. Depends on the SEO content pipeline.
- **Video as a view-through asset.** The growth-team video runs as a view-through demand asset feeding retargeting into Search and LinkedIn, not just a single conversion campaign. Depends on the video pipeline.
- **A second incrementality holdout.** After native-form leads are traced to pipeline, run a Meta value holdout (the brand-defence geo holdout comes first, at launch). We earn measurement sophistication as volume grows rather than front-loading it onto a rebuilding account.

Appendix A. ABM build detail

Short version: account-based done the way it actually works. Broad-but-precise audiences, native capture, nurture behind the click. This folds in the former standalone ABM spec so it all lives in one place.

The old ABM was not a bad idea executed unluckily. It broke three rules at once, and the audit shows each break costing us: narrow hand-picked Top-100 lists that sat below LinkedIn's delivery threshold, pushed through InMail, with no nurture layer behind the click. The one broad, correctly-targeted campaign we ran did the opposite. That contrast is the whole argument for the rebuild.

SAME OFFER, SAME FORM, TWO AUDIENCES	AUDIENCE	SPEND	LEADS	COST PER CLICK
Narrow "Top 100" list (about half the wrong vertical)	3,400	€441	0	€184
Broad ICP audience	22,000	€444	3 (21 lifetime)	€8.28

Near-identical spend. The only variable was the audience. Narrow targeting starved the platform of the signal it needs to deliver, and cost 22 times more per click for zero leads. Figures as recorded in the audit.

Two more lessons sit underneath that. First, native lead forms beat landing pages by a wide margin: on Meta the same ebook offer captured leads at €26 through a native form against €84 through a landing page, verified from the platform export. Second, none of it matters if the lead then hears nothing, because most account-based targets aren't ready to buy on first touch. The rebuild fixes all three: broad-but-precise audiences, native capture, and a nurture layer that wasn't there before. The position to hold internally is simple: we aren't killing ABM, we're fixing the mechanics that made it fail.

The three-tier model (recap)

The rebuild replaces the single narrow list with three tiers that do different jobs (the same tiers used in section 6). Most of the work and most of the pipeline come from the middle tier.

TIER	WHO IT TARGETS	AUDIENCE	TARGETING METHOD	FORMAT	PRIMARY METRIC
Tier 1: Named accounts	A small, hand-chosen set of priority logos	10-30 logos	Company-list, one-to-one creative	Conversation ads, native forms, sales-aligned	Meetings booked
Tier 2: Best-fit (the workhorse)	Verticals and sizes that match the product-market map	50K+	Title + industry + size, above the floor	Native lead forms	Cost per lead, then pipeline
Tier 3: Broad ICP + lookalikes	The wider ICP and lookalikes of converters	Largest	Broad ICP + lookalikes from Tier 2 and closed-won	Native forms, retargeting	Lowest cost, feeds the funnel

The tiers are a funnel, not silos. Tier 3 finds and warms people cheaply, the best graduate into Tier 2, the genuinely strategic accounts get the Tier 1 one-to-one treatment. Closed-won data flows back to sharpen the Tier 3 lookalikes, so the system gets better as it runs.

Build standards every ABM campaign must meet

Each standard exists because the audit caught the opposite failing somewhere. A campaign doesn't go live until it meets all of them.

STANDARD	THE RULE	WHY (FROM THE AUDIT)
Audience floor	50,000+ for any always-on tier. No hand-built lists below the delivery threshold.	The 3,400 narrow list spent €441 for zero leads. The 22K broad audience worked.
Native capture	Native lead forms by default. Landing pages only with a specific reason and a working pixel.	Native €26 per lead vs landing page €84 on the same offer.
Conversion event wired	Lead MQL and Lead SQL confirmed firing before launch, not after.	The conversion signal was broken across platforms, so spend couldn't be judged.
Right bid strategy	Conversion-based bidding into that event. Never Maximize Clicks on an always-on campaign.	Maximize Clicks bought volume at near-zero conversions across the biggest campaigns.
Suppression and exclusions	Suppress existing customers, open opportunities and clearly wrong verticals from day one.	Half the narrow Top-100 list was the wrong vertical. Budget chased people who couldn't buy.
Naming convention	Dual-string: the Dynamics campaign theme as prefix, PPC detail as suffix, separated by an underscore.	19 of 120 paid opps arrived with no campaign attached. Consistent naming closes that gap.
Nurture behind the click	Every lead enters a sequence by segment before any sales hand-off. No lead goes cold in silence.	Paid generated leads but had no nurture layer, so most of what it captured went cold.
Proven creative	Start from the hooks that worked (the training-to-execution angle, the Easilys video), not from scratch.	The creative was rarely the problem. Targeting, budget and capture were.

What a campaign looks like when it is built right

Three example builds, one per tier, for the French Easilys line. These are planning targets to build against, not results. The unused 1.28-million-member French matched audience we already hold is the starting pool for Tiers 2 and 3.

TIER 1 EXAMPLE	FR EASILYS, NAMED ACCOUNTS
Audience	10-20 priority contract-catering and public-sector logos, hand-chosen with sales
Format and bid	Conversation ads plus native form, conversion bidding, sales sequencing in parallel (not paid)
Budget bucket	Evergreen (always-on, strategic)
Target	Meetings booked, not cost per lead. A handful of the right meetings is the win
Depends on	Tracking fix for account-level scoring and pipeline attribution

TIER 2 EXAMPLE (THE WORKHORSE)	FR EASILYS, BEST-FIT SEGMENTS
Audience	Contract caterers, hotel groups and FSR chains, filtered from the 1.28M pool by title, industry and size, kept above 50K
Format and bid	Native lead form, conversion bidding into Lead MQL, suppress customers and open opps
Budget bucket	Evergreen (always-on demand generation)
Target	A defensible cost per lead first, then lead-to-MQL and MQL-to-pipeline once tracking is live
Build now?	Yes. Segment and sequence design has no dependency. Audience sync waits on the tracking fix

TIER 3 EXAMPLE	FR BROAD ICP PLUS LOOKALIKES
Audience	Wider French hospitality ICP, plus lookalikes built from Tier 2 leads and closed-won
Format and bid	Native forms and retargeting, lowest-cost reach, feeds the Tier 2 pool and the remarketing audience
Budget bucket	Evergreen (always-on)
Target	Lowest cost per lead in the account, and a steady supply of warmed contacts graduating upward
Build now?	Lookalikes wait on closed-won feeding back through the tracking fix. Broad ICP can start sooner

Where ABM sits in the budget. The always-on ABM tiers sit in the evergreen budget rather than the growth-support share. Tiers 2 and 3 run year-round capturing and feeding demand, which is the definition of evergreen, not event support. That frees the growth-support share for the growth team's own campaigns.