

Paid Media and Lifecycle: Audit, Findings, and Forward Strategy

Adrian Fernandez | 30 May 2026, revised 2 June 2026 | Working document

Executive summary

This report reviews the paid media operation we inherited across Google, LinkedIn, Microsoft and Meta in the UK, Iberia, France and LATAM, and sets out how paid media and lifecycle nurture become one engine rather than two separate efforts.

Two findings frame everything that follows. The first is that paid is a stronger source of business than the day-to-day reporting suggests. Over the last twelve months it produced 120 marketing-sourced opportunities and 19 closed-won deals, which places it second only to the demo-request bucket on won business, ahead of website, events and email. The second is that almost everything holding paid back is fixable, and most of it without new budget. The audit found spend set for clicks rather than customers, a conversion signal that was not firing into the platforms, brand defence switched off for five months, and no nurture layer to convert the leads that were not ready to buy yet.

The recovery from cutting the identified waste is roughly 35,000 to 55,000 euros a year, on top of about 9,000 already saved by stopping a failed campaign. No new budget is required: the reactivations and the rebuild are funded by redeploying what is currently wasted.

One principle governs how we judge paid from here: we measure it on incremental lift, proven by holdout, not on platform-reported conversions, and we move budget to where the lift is real.

The sections below cover where paid stands among our sources, the findings and their root causes, the platform-by-platform detail, the money, the forward plan, and the expansion motion that grows the customers we already have.


How to read this, and where the numbers come from

The windows differ by necessity. Google and Microsoft figures come from a 90-day audit window, annualised where stated. Meta figures come from a 17-month export, because that is the period over which the account ran. The opportunity, pipeline and won-deal figures come from the last twelve months in Dynamics, the system in which sales actually work deals, covering 875 marketing-sourced opportunities created between 28 May 2025 and 28 May 2026.

One note on the source figures, because it matters later. A separate filtered view of the same Dynamics data, shared earlier, shows smaller totals and ranks paid third by pipeline value. That view and the full trailing-twelve-month export do not reconcile on the number of opportunities they include, which points to a different window or filter behind it. This report uses the full export, which can be reproduced line by line. Confirming the basis of the filtered view is an open item, noted again at the end.

1. Where paid stands among our sources

Grouped by pipeline value, paid sits fourth among marketing sources over the last twelve months. That is the least flattering way to read it, and it is worth stating plainly before making the case that it understates paid.




Request a Demo is a last-touch bucket, not an acquisition channel: it holds demand other channels created.

Two things make the value ranking misleading. The first is the bucket sitting on top of it. "Request a Demo" is not an acquisition channel, it is a last-touch destination: it records where an opportunity was sitting when it converted, not what created the demand. A large share of those 394 opportunities are paid, event and website visitors who finally filled in a demo form. Counting it as a source flatters it and starves the channels that actually fed it. The second is that pipeline value counts opportunities created, not business won, and on that measure the order changes sharply.

When you look at deals actually closed, paid is second only to the demo-request bucket, and it is comfortably ahead of website, events and email.

Closed-won deals by source (trailing 12 months)



Paid is second only to the demo-request bucket on closed business, ahead of website, events and email.

The clearest way to see why is win rate, the share of each source's opportunities that convert to won business.

SOURCE	OPPORTUNITIES	PIPELINE VALUE	WON DEALS	WIN RATE
Request a Demo <i>(last-touch bucket)</i>	394	€6.04M	62	15.7%
Marketing Event	138	€1.73M	7	5.1%
Marketing Email	120	€1.37M	2	1.7%
Paid (search + social)	120	€1.05M	19	15.8%
Website	81	€0.66M	13	16.0%

Marketing-sourced opportunities, Dynamics, trailing 12 months. Win rate is won deals divided by opportunities. Red marks the sources that carry pipeline but rarely convert it.

Paid converts at 15.8 percent, essentially the same as website (16.0) and the demo bucket (15.7), and far above events (5.1) and email (1.7). Marketing email is the sharpest contrast: it carries more pipeline value than paid, 1.37 million euros against 1.05 million, but it has produced only 2 won deals against paid's 19. So the channels ranked above paid on pipeline value are, with the exception of the last-touch demo bucket, weaker than paid at turning that pipeline into revenue. On the measure that pays the bills, paid is a top-two source.

Why paid still looks small in the day-to-day reports

If paid is a top-two source on won business, why does it read as minor in HubSpot's standard source reports? Because of three breaks in the reporting layer, all of them upstream of the HubSpot attribution model, which is itself sound.

The first is legacy option IDs. Historic deals store their source under IDs from an earlier version of the source dropdown. The label still displays correctly when you open a record, so it looks fine, but reports that group by source match on the current IDs and read those older deals as blank. The second is that deal source is written by the Dynamics sync rather than by tracked web activity, so without consistent UTMs in place a genuine paid lead can arrive in HubSpot with no campaign attached. The third is last-touch attribution: a paid click that converts weeks later on a demo form is credited to the demo, not to paid, which is the same effect that inflates the demo-request bucket.

None of these is a flaw in how attribution was designed. They are a reorganised picklist and a sync behaviour sitting in front of a sound model. The fix is a one-time re-map that reads each deal's correct label and rewrites it to the current option ID, which would recover almost the entire history. This needs verifying on a sample of deals before it is relied on, by checking the property change history, and it sits with RevOps with marketing sponsoring it. Roughly 10,900 main-pipeline deals currently read blank for this reason, and about 25,000 across all pipelines, so the upside of fixing it is large.

There is a fourth effect, and it is operational rather than a reporting break. Today only the Request a Demo form auto-routes a lead to an SDR. Every other form a paid lead can fill, a whitepaper or guide download, an event registration, a product-page enquiry, creates an event in HubSpot and then nothing happens to it. So a real share of the demand paid pays to capture is never worked and never progresses, which means it never shows up as paid-sourced pipeline either. That gap is addressed in the forward plan, and it is part of why paid reads smaller than it is.

2. The seven findings, and the three causes behind them

The audit produced seven findings that matter. They are easier to act on grouped by the three root causes they trace back to, because fixing the causes stops the individual problems recurring.

Money pointed at the wrong target. Bidding across the top-spending campaigns was set to Maximize Clicks, which buys the most clicks for the budget regardless of whether those clicks convert, and on Meta 73 percent of spend ran on Traffic and Awareness objectives rather than lead generation. Budget also went to products a market cannot sell: Workforce campaigns running in France, and a LATAM Workforce campaign nested inside the Spain account, both against the actual product-market map. This is findings two and four.

Results were invisible. The lead conversion events were not firing reliably into the ad platforms, and a main Meta pixel sat dead for about four weeks. The root cause sits upstream of the platforms and of attribution: the website migration changed the dataLayer contract those conversions depend on, so the conversion actions are live but stopped being triggered. When the conversion signal is broken, good and bad spend look identical, the platforms cannot steer toward conversions even if asked to,

and nobody can see which campaigns are working. Alongside this, source attribution was under-reporting paid for the reasons in the previous section. This is findings one and six.

Demand was wasted after capture. Paid generated leads, but there was no nurture layer to work the ones that were not ready to buy on first contact. In B2B SaaS most buyers are not, so a lead that fills a form and then hears nothing goes cold. This is finding seven, and it is the one that matters most for the forward plan, because it is the difference between paid as a lead source and paid as a pipeline source.


Two further findings are execution misses that the rebuild corrects directly rather than root causes. Brand defence on Mapal's own branded search was paused on both Google and Microsoft in the UK for five months, leaving the cheapest converting campaign in the account, 13.64 euros per conversion across 162 conversions, switched off while competitors bid on the brand unopposed. And narrow, hand-built audiences consistently underperformed broad, correctly-targeted ones: on LinkedIn the same French offer through the same lead form produced zero leads from a narrow 3,400-person list that was half the wrong vertical, against three leads at 148 euros each from a broad 22,000-person audience, and native lead forms came in roughly three times cheaper than the same offer sent to a landing page.

3. The platform audit

The same pattern, click volume on a broken conversion signal, shows up across all four platforms, but the specifics and the fixes differ by platform and market.

Google Ads

Google is the largest spend and the clearest case. The year-on-year picture is the headline: across the three markets, cost rose sharply while conversions fell.



Cost up across every market, conversions down where they can be measured. France conversions are not comparable year-on-year.

In Spain and LATAM, nine of the ten top-spending campaigns ran on Maximize Clicks. The Workforce-ES search campaign spent 7,305 euros for a single conversion, with 6,593 of that on mobile for none, while a defunct "Compra" conversion was still set as the primary goal so the bidder was partly aiming at an event that no longer exists. The fix is to move the flagship onto conversion bidding into a working Lead MQL event, cap the secondary products, and either fence or switch off Display, which was running with no placement exclusions. The negative-keyword hygiene was the one bright spot, four product-segmented lists in active use.

In France, the product focus was right, Easilys is the correct French flagship, but the execution diluted it across twelve overlapping "Scale" variants that split the budget and the conversion signal twelve ways. Roughly 9,160 euros, 44 percent of the French Google budget, ran across 52 keywords that were off the core play. The conversion layer itself was broken, with 18 of 29 actions in a Removed state and four removed actions still flagged as primary. The fix is to consolidate to one search campaign and one Performance Max, both feeding a single clean conversion. Performance Max already converts when the signal reaches it, which tells us the offer works.

In the UK, generic terms ran on Maximize Clicks at zero conversions, the Flow training campaign spent 3,445 euros for none, while brand defence, the cheapest converter in the account, sat paused since early December. There are no negative-keyword lists at all and roughly 69,000 placement exclusions attached to nothing. The single biggest lever here is simply turning brand defence back on.

LinkedIn Ads

LinkedIn ran narrow account lists below the platform's delivery threshold, and the controlled comparison described in the findings, narrow versus broad on the same offer, is the clearest evidence in the whole audit that broad, correctly-targeted audiences win. A ready 1.28-million-member French matched audience sat unused while the narrow wrong-vertical list ran. The infrastructure is healthy, the Insight Tags fire on all three accounts, and the UK training-to-execution creative is reusable. The fix is broad ICP audiences above the 50,000 floor with native lead forms, and leaving narrow account targeting to sales sequencing rather than paid.

Microsoft Ads

Microsoft is effectively a new channel rather than a fix. The UK account is entirely paused, roughly 393 euros a day of capacity


sitting idle, and it was configured with a 6-euro target cost per acquisition on software that sells for 100 to 400 euros, a Helsinki time zone, and euros instead of pounds. The French side is 42 campaigns auto-imported from Google on Maximize Clicks, with the Audience Network alone taking 31 percent of French Microsoft spend for almost nothing, and the best curated negative-keyword list in the account, 250 keywords, attached to zero campaigns. The French Easyls brand term does convert at around 143 euros. The right framing is to decide what Microsoft should be, brand defence plus high-intent, and to open a Spain account, which does not exist today.

Meta Ads

Meta ran for 17 months and spent 14,958 euros for about 108 leads, with 73 percent of that, roughly 10,857 euros, on Traffic and Awareness objectives rather than leads. The Audience Network leaked 3,421 euros, the main pixel was dead for about 27 days, and four other pixels are inactive. The one thing that worked is the native lead form, which captured leads at roughly three times the efficiency of the landing-page route, 26 euros against 84. That efficiency has not yet been traced through to closed pipeline, and the form cost regressed from 26 euros in its first version to 55 in its second, so the next step is to trace those 108 leads to deals and diagnose the regression before scaling. France stays the live base.


4. The waste, and what we get back

The audit identified roughly 26,000 euros of waste over the 90-day window, which annualises to about 105,000 euros if the pattern continues. The largest items are concentrated in a handful of campaigns.



About €26,000 over 90 days, roughly €105,000 a year at this rate. Red is pure waste, amber is wrong placement.

Applying confidence haircuts, higher for clear-cut waste and lower for items that need confirmation, gives a realistic recovery of 67,410 euros a year. The number to use externally is more conservative still, to allow for execution risk and cost-per-click inflation.



No new budget. The reactivations and the rebuild are funded by redeploying what is currently wasted.

The full line-item detail, with the haircuts applied, is below.

ITEM	PLATFORM / MARKET	90-DAY	12-MO RUN-RATE	RECOVERY	CONFIDENCE
Workforce-ES search, mobile	Google / Spain	€6,593	€26,755	€22,742	High
France keyword spread (12 variants)	Google / France	€9,160	€37,160	€20,438	Medium
Hospitality Jobs audience (mismatch)	Google / UK + FR	€5,356	€21,737	€11,955	Medium
Spain Display (4 campaigns)	Google / Spain	€3,156	€12,806	€7,043	Medium
Microsoft Audience Network	Microsoft / France	€901	€3,656	€2,011	Medium
Microsoft Workforce (4, off-map)	Microsoft / France	€785	€3,185	€2,707	High
Easily's without negative list	Microsoft / France	n/a	n/a	€400	Medium
Competitor brand bids	Microsoft / UK	€94	€381	€114	Low
Total realistic recovery		€26,045	€105,680	€67,410	

Two honesty notes. Meta's Audience Network leak of 3,421 euros is a 17-month figure on a different clock, so it is kept out of the 90-day total rather than stacked on it. And the larger Meta number, the roughly 10,857 euros on the wrong objectives, is not counted as waste at all: it bought real clicks, it was simply pointed at the wrong goal, so it is budget to redeploy rather than money lost.

5. The forward engine: paid and nurture as one loop

The single most valuable change is not more spend, it is building the nurture and remarketing layer that did not exist before, because it works demand we have already paid to create.

The loop is straightforward. Paid search and lead forms capture the lead. HubSpot segments it by product and market and by what brought it in. A nurture sequence works the leads that are not yet ready, while paid remarketing keeps the brand present between emails so the lead does not go cold. When the lead shows real intent, pricing or demo-page activity, its score crosses a threshold and it is routed to sales as a warm hand-off rather than a cold one. The outcome, won or lost, feeds back into the audiences and the bidding so the system learns from actual revenue.

The behavioural segments and the nurture sequences can be built now, because they have no dependency on the tracking fix. Syncing those segments out to the ad platforms for remarketing, and building revenue-based lookalikes, comes once the conversion tracking is healthy. The stages and what triggers each one are below.

STAGE	TRIGGER	WHAT RUNS	GOAL
New paid lead	Form fill from a paid click, not a demo	Tag the source, send a welcome and relevant content by product and market	Establish relevance fast
Engaged, not ready	Content downloads, repeat visits	Sequenced nurture by segment, branched on what brought them in	Move toward sales-ready
Showing intent	Pricing or demo-page views, email clicks	Lead score crosses threshold, priority routing, sales alerted	Hand over warm, not cold
Active in nurture	In any segment	Paid remarketing stays present across Google, LinkedIn, Meta	Stop the lead going cold between emails
Closed or disqualified	Sales outcome recorded	Outcome feeds back to audiences and bidding	The loop learns from real revenue

The routing gap this loop had to close, now largely closed

The loop only works if a captured lead actually moves, and until recently most non-demo leads did not: only the demo form auto-routed to an SDR, and every other form a paid lead fills, a whitepaper or guide download, an event registration, a product-page enquiry, created an event in HubSpot and then sat. That has now changed. A central routing workflow is live that creates, qualifies and routes the non-demo leads to an owner, splits current customers to Customer Success, suppresses duplicates against any open lead, and blocks the clear non-fits such as non-hospitality and very small companies. For paid that closes the worst dead-end: the leads we pay to capture are now created and routed rather than lost.

Two gaps remain, and both are small. SDRs are alerted by email only for demo requests, so a routed content or event lead can still sit if nobody is prompted to work it; the fix is a light working signal on the non-demo paid leads, a shared queue view or a daily digest rather than a per-lead alarm, raised with RevOps and the SDR side. And paid leads need a clean paid-source label on entry so they are visible and reportable through the routed flow, which is ours to set. The table below is the current split.

PIECE	STATUS	OWNER
Create, qualify and route non-demo leads to an owner	Live. Customers split to Customer Success, duplicates and clear non-fits filtered out	RevOps (shipped)
Paid-source tagging on entry	To set, so paid leads are visible and reportable in the routed flow	Us
Working signal on non-demo leads	Residual gap: only demo requests alert an SDR today; a queue view or daily digest closes it	RevOps + SDR side
Nurture sequences by segment	No dependency, build now; works the leads between touches	Us
Automatic scoring and graduation	Grades leads on fit and intent rather than a single action; swaps in cleanly when it lands	RevOps, Phase 2

Graduation can still route two ways. The default is the SDR queue; for named or higher-value accounts, or a market where we choose a softer first touch, a lead can route instead to a LinkedIn-connect task for the right salesperson, tracked the same way so we can compare which path converts. The remarketing sync waits on the conversion-tracking fix; the nurture itself works without it, so it starts now.

6. Growing the customers we already have

The forward engine in section 5 turns strangers into customers. This section extends it to the customers we already have, the third loop the plan carves out and points to whenever the acquisition audiences suppress current customers. Acquisition wins the logo; expansion grows it. In a multi-module business the cheapest pipeline we have is the base we already serve, because the relationship and the data already exist and existing customers convert far better than cold prospects.

Two motions. Cross-sell is a customer on one module taking an adjacent one. Up-sell is growing what they already have: more sites across the estate, more seats, a higher tier, a longer term. Both are measured as expansion revenue and net revenue retention, not new-logo pipeline, and neither competes with acquisition for budget because the audiences are the customers we already hold.

The cross-sell map. It follows the same product-market map as acquisition, with the same product, market and ICP personalisation, so a Spanish Workforce customer gets a Spanish Compliance cross-sell carrying Spanish proof.

CURRENT PRODUCT (ENTRY)	NATURAL NEXT MODULE	THE ANGLE
Flow (UK, hotels)	Compliance, then Engage	Close the compliance and audit gap on teams already trained. Flow and Compliance are often bought together.
Workforce (Iberia)	Compliance + Engage	The compliance and engagement layer on the same teams whose scheduling and labour cost are already solved.
Easilys (France)	Flow	Bring training and onboarding into a kitchen operation that already runs on Easilys.
Key Accounts	MapalOS consolidation	More modules on one operating system, one login, one data layer. The largest-account expansion play.

How it runs. Marketing owns the demand layer: a customer-only nurture stream, renewal sequences at 90, 60 and 30 days, and customer-match ads that surface the right module at the right moment. It then hands a qualified expansion signal to the account owner. Customer Success and account management own the sale. The customer-to-Customer-Success routing and the marketing-attribution capture on customer deals already exist in the live routing workflow, so the handoff path is in place; what the expansion build adds is the customer nurture, the health gate, and the paid customer-match layer. One hard rule sits above everything: only green-health accounts get an expansion ask, and an at-risk account is a save rather than a sell. This is a post-launch Q3 build, gated on two data feeds beyond the acquisition tracking work, both targeted Q3: product-usage into the CRM (owner: RevOps) and the account-health score exposed to marketing (owner: Customer Success), and it is not part of the paid launch.

How we measure it. Expansion revenue and net revenue retention, the share of customers on two or more modules, estate roll-out and renewal uplift, and the expansion opportunities the customer programme influences. It is judged on incremental lift, proven against a held-out set of matched green-health accounts, the same discipline section 8 sets for acquisition spend. The aim is a third pipeline source standing beside new-logo paid and the nurture programme, built almost entirely on demand we have already paid to create.

7. The rebuild, the budget, and the ABM tiers

The rebuild concentrates the current sprawl into roughly 33 campaigns, about four on Google, three on LinkedIn, two on Microsoft and three on Meta per active market, each funded above the platform's learning threshold, each with one flagship product per market, and each measured on a working conversion. One person can run weekly hygiene across 33 focused campaigns. The current spread cannot be managed properly by anyone.

No new budget is required. The indicative monthly allocation under the new shape is below, in planning ranges to be firmed up once reliable cost-per-lead by channel is in place.

PLATFORM	UK	IBERIA	FRANCE	LATAM	MONTHLY
Google	€4-6K	€4-6K	€4-6K	€0.6-1K	~€14-19K
LinkedIn	€3-4K	€3-4K	€1.5-2K	-	~€7.5-10K
Microsoft	€1-2K	Q3 build	€0.8-1.6K	incl.	~€2-3.5K
Meta	€0.5-1K	€0.5-1K	€1-1.5K	-	~€2-3.5K

Planning ranges, not commitments. Funded by redeploying current waste, not by new budget.

Within each platform, budget splits between always-on evergreen activity and a ring-fenced growth-support share for the growth team's events, webinars and content. Search stays entirely evergreen, because intent capture should never be sacrificed to promote an activity. All always-on demand generation lives in evergreen, including the account-based tiers below, which run year-round. The growth-support share funds the growth team's own campaigns, and only an account-based burst built specifically to amplify a growth-team event or content drop draws on it. The split is roughly 70/30 on LinkedIn and 75/25 on Meta, with Google and Microsoft fully evergreen, and the growth share is a ceiling, not a floor: if it is not used it reverts to evergreen.

Account-based marketing is rebuilt in three tiers rather than as the narrow lists that underperformed. The top tier is a small set of named accounts, 10 to 30 logos, run one-to-one and measured on meetings booked. The middle tier, the workhorse, is best-fit segments built into audiences above the 50,000 threshold and measured on cost per lead through to pipeline. The base tier is broad ICP plus lookalikes, feeding the top of the funnel at the lowest cost. The earlier account-based work underperformed on mechanics, tiny lists, InMail and no nurture, not on the idea, and this keeps the intent while fixing the mechanics. Because all three tiers run year-round, they sit in the evergreen budget rather than drawing on the growth-support share, which frees that share for the growth team, and the unused 1.28-million-member French audience becomes the base.

8. Competition, sequence, and how we measure it

The competitive read changes where budget should point, because the set we actually compete against is not the one assumed in HQ thinking. In France, Lightspeed tops the auctions at around 52 percent, with Skello consistently present: it is a point-of-sale and Easilys fight, which is the argument for keeping Workforce out of the French paid mix. In Iberia, the live competitors are HR and payroll tools such as factorial and tramitapp rather than point-of-sale, so Workforce stays but the messaging competes on the HR axis. In the UK, the set skews to LMS and training tools such as Absorb, which puts Flow's training-to-execution angle on the right axis. The workforce-management names HQ assumes, such as Fourth and Deputy, barely appear in any auction, so we plan against the competition we can actually see, and brand defence matters most where rivals bid on our terms, which is confirmed on French Bing.

The work is not blocked. Native lead forms can capture leads now, and the dependencies govern measurement and smarter bidding rather than whether we can run at all. The 90-day sequence runs from stabilising in the first two weeks, to replatforming bid strategy as tracking lands, to replicating the lead-form model and standing up nurture, to building out the newer surfaces.

PHASE	DAYS	WHAT HAPPENS
Stabilise	0-14	Brand defence back on, kill the clear waste, fix conversion goals, Audience Network exclusions
Replatform	15-30	Move from Maximize Clicks to conversion bidding as tracking lands, consolidate Easilys
Replicate	31-60	Lead-form model to UK and Spain, first nurture sequences live, audiences syncing
Build	61-90	Spain Microsoft account, France landing page, remarketing on populated audiences

Four things have to happen, in order: the conversion-signal fix, which is the critical path the bidding switch waits on and which sits with us and the website-migration lead, because the campaign conversions are live but the site migration changed the dataLayer contract that triggers them, so the fix is in the tag container and on the site, not in attribution; the cross-channel UTM standard, in progress; the legacy source re-map and the Dynamics source-overwrite, which sit with RevOps and need verifying on sample deals first; and leadership sign-off on consolidating the campaign count, because cutting campaigns reads as doing less without the rationale behind it.

One operational constraint shapes how the nurture-to-SDR routing is enforced. The latest pipeline review shows SDR outbound running at about 63 percent of target, roughly 87,000 against 138,000 euros, on small weekly targets and real load. A rigid speed-to-lead rule, call within 30 minutes or escalate, on a team already behind would create noise and resentment and get ignored. So paid leads get queue priority and a daily monitoring report with soft escalation, not a punitive timer. It moves behaviour where it matters, protects the relationship with a stretched team, and can be tightened to a hard service level later if the soft version is not enough.

Reporting steps up as each measure comes online. Weekly: spend, leads and cost per lead by platform and market. Monthly: MQL and SQL volume, MQL-to-SQL conversion, and spend efficiency. Quarterly: pipeline contribution, ARR, expansion revenue and net revenue retention, and a full re-baseline against the 67,410-euro projection, run 30 days after the tracking fix lands so it is measured against real numbers. The target is for paid to hold its place among the top sources of pipeline and ARR, beside website, SEO and the nurture programme it now feeds. The floor is already real: 19 won deals and second place on closed business, achieved while half the measurement was broken.

One principle now governs how we judge paid going forward, and it sharpens the brand-defence case rather than weakening it. Platform-reported conversions over-count, because a paid click that would have happened anyway, for free, is not new business.

So we reactivate brand defence, but we treat its value as a hypothesis to test rather than an article of faith. The strongest version of that hypothesis is where competitors bid on our terms, which is confirmed on French Bing: there, a branded click we would otherwise lose to a rival is genuinely incremental, and defending it is worth real money. We prove it with a geo holdout, brand defence switched on where rivals bid our brand and off in a comparable market where they do not, comparing branded conversions over four to six weeks. If the conversions hold steady with paid off, the spend was not incremental and the budget moves to where lift is real. The standard this sets applies to every channel: we judge paid on incremental lift, measured by holdout, not on platform-reported conversions. The discipline comes from large-scale holdout testing in the industry, most famously the Uber case, where a single set of experiments reallocated roughly 35 million dollars once it showed how much cheap last-click spend was buying conversions that would have arrived anyway. We would rather move budget to where the lift is real than spend because the budget exists.

Open items

Two things remain to confirm before this goes outside the team. The first is the basis of the filtered source view that shows paid third by value and smaller totals, so its window and filter can be reconciled with the full export this report uses. The second is tracing Meta's roughly 108 leads through to closed pipeline, which is what would turn the native-lead-form efficiency from a cost story into a revenue one. Neither changes the central finding: on closed business over the last twelve months, paid is a top-two source, and it has been under-credited rather than underperforming.

One option worth holding, surfaced by the attribution work: because Dynamics holds clean source for every channel, not just paid, a full-funnel marketing-source view could be built straight from Dynamics without waiting on the HubSpot connector. It is not part of this plan, but it is the most useful thing the attribution review turned up.